Stories Coffee — Recommendations

Based on 51,177 item lines and 33,887 baskets across Garden, Harbour and Uptown, 1 Jan – 30 Jun 2026 ▪ $305,925 revenue ▪ $179,430 recorded gross profit ▪ 17 October 2026

Read this first. The file looks simulated. Transaction IDs start with `SYN-`. Harbour's prices are exactly Garden's $\times 1.05$ and Uptown's $\times 0.96$. Prices and costs never change in 181 days. Hourly demand is pure random noise around a fixed weekly schedule. The March 2026 war, when Lebanon's restaurant syndicate reported falls of over 85%, does not appear at all. Every recommendation below is therefore a *tested pilot*, not a rollout. Reconcile the till with card settlements and cash deposits before committing capital.

At a glance

#	Recommendation	Size of the prize / signal	Confidence
1	Audit Garden's costs before changing suppliers	$3,981 – $15,923 GP per 6 months	High (gap), cause open
2	Staff, prep and seat by the clock, not the weather	Template within 1.2% of best possible forecast	High
3	Move the 10% discount out of peak hours into quiet ones	Same budget; mornings -37%, weekends -29%	Medium (test)
4	Stop 8-day blanket 15% campaigns; time them to slow periods	No measurable daily GP lift	Medium–high
5	Run a no-discount evening attachment pilot: Cold Coffee + Bakery	Evening = 44.6% of GP	Medium (test)
6	Test a uniform +5% price at one branch for four weeks	≈$7.3k GP per 6 months if demand holds	Medium (test)
7	Test later opening at the weakest early hours	Median $656 per 6 months, pays in 75% of draws	Low–medium

1. Audit Garden's costing rule before touching procurement

Garden records a **46.91%** gross margin, against Harbour's **64.96%**, on the same 45 products.

- Garden's unit cost is $\approx 1.43\times$ Harbour's on every product, and it never changes. That looks like an allocation rule, not waste.
- Repricing Garden's quantities at Harbour's costs adds **$15,923.47** in six-month GP. If only 25–100% of that is real cash, the range is **$3,981 – $15,923**. If it is purely an accounting allocation, the saving is $0.

Do this: Trace the largest dollar contributors (the top five hold 22.3% of the gap; Iced Matcha and the Chicken Pesto Sandwich lead) from the POS cost master through recipes to supplier invoices.
Go / no-go: change suppliers or recipes only if invoices show an avoidable cash difference larger than the cost of acting. At 50% recovery the saving survives up to 14.5% volume loss, so protect quality. **Confidence: high that the gap exists; the cause is unknown.**

2. Run operations off an hour × weekday template, not the weather

Across all 81 features we engineered (weather, dust, holidays, Ramadan, school breaks, customer history), only the clock survives controls.

- A branch × hour × weekday template predicts unseen May–June hours with an error of 1.59 baskets. A perfect model would still miss by 1.57, so the template is within **1.2%** of the limit. Adding weather changes nothing (1.594).
- 19:00–24:00 brings **44.2%** of baskets and **44.6%** of GP. Dine-in jumps from **46% to 72%** at 18:00, and cold-drink share rises to **70%** after 19:00, in every month and branch.
- The busiest slot is Harbour on Tuesday at 21:00: 21.3 expected baskets in June, **plan for 27** (90th percentile). Tuesday runs about $1.5\times$ other weekdays.

Do this: roster staff, ice, cold-brew prep, desserts and seating from the staffing template (in the project database). Log queues and stock-outs from 19:00 to 23:00 for four weeks.
Go / no-go: add evening labour only if logged lost orders × $5.34 GP per basket exceed its cost. **Confidence: high.**

3. Spend the existing 10% discount where demand is weak

Today the discount ignores the clock. About **5.5%** of items are discounted in the morning, afternoon and evening alike, and on every weekday (5.1–6.1%).

- Mornings bring **37%** fewer baskets per hour and weekends **29%** fewer per day. Yet over half of discounted items (52%) are sold in the busy evening, to customers who would come anyway.
- The payday window (last 3 and first 3 days of the month) already brings +4.7% daily baskets, so don't discount into it.

Do this: for four weeks, at one branch, give the 10% offer only before 11:00 and on Saturdays. Keep the other branches as controls.
Go / no-go: keep it only if GP per assigned visitor rises against control. Basket averages mislead, because bigger baskets are more likely to contain a discounted item by chance. **Confidence: medium.**

4. Retire blanket 8-day 15% campaigns; rerun them as a timed experiment

The two 15% windows, 12–19 March and 12–19 May, ran at Harbour on Cold Coffee and Desserts across every weekday and at peak hours.

- Against Garden and Uptown over the same dates, Harbour's daily GP during the campaigns was *lower*, with wide intervals. Against Harbour's own no-discount baskets, the 15% baskets earned **$0.33** less GP.

Do this: randomize 0% / 10% / 15% offers by day at Harbour and log people who don't buy.
Go / no-go: scale only an arm that raises contribution per assigned visitor after discount cost. **Confidence: medium–high that the current format doesn't pay.**

5. Pilot evening attachment without a discount: Cold Coffee + Bakery

The evening is a sit-down cold-drink and dessert occasion: Cold Coffee is 16.6% of evening units, Matcha 17.5% and Desserts 10.7%. Dine-in and takeaway baskets earn the same GP ($5.32 vs $5.38), so the gain must come from more items per basket, not from seating.

- Cold Coffee + Bakery repeats in both halves of the period, carries a 66% recorded margin, and coffee input costs are falling. Matcha supply is being rationed and its prices set records, so don't stimulate demand for it before securing supply.

Do this: a scripted "add a pastry?" prompt from 19:00 to 23:00 with no discount, run on alternating weeks.
Go / no-go: keep it if attached GP exceeds labour time and any service slowdown at the peak. **Confidence: medium.**

6. Test a uniform +5% price rise at one branch

At Stories' 58.7% margin, a 5% rise pays unless demand falls by more than 1.68% for every 1% of price. A 20,000-draw simulation puts the gain at **$7,305** per six months (5th–95th percentile $3,025 – $11,128). That certainty comes from the assumptions, not from data: prices never changed in the file.

- Holding coffee prices back loses to a uniform rise in 88% of draws. Eleven Harbour prices sit a few cents above a whole dollar, the worst place for a price.

Do this: +5% at one branch for four weeks against the other two, with price endings cleaned up.
Go / no-go: stop if baskets per day fall by more than the break-even drop of about 8% over the test period. **Confidence: medium.**

7. Test later opening at the weakest early hours

Uptown's 07:00 hour averages $4.37 GP against an estimated $6.53 in staff and generator cost. The generator tariff rose 74% from January to April. Closing Uptown 07:00–09:00 and Garden 07:00 pays in 75% of simulated scenarios, with a median of $656 per six months.

- Staffing assumptions drive this result far more than energy prices. Confirm the real early-shift roster first.

Do this: open Uptown at 09:00 and Garden at 08:00 for four weeks.
Go / no-go: reverse it if those customers don't come back later in the day. **Confidence: low–medium.**

What we would *not* do — and why

- **End-of-day discounts.** Late evening is Stories' *peak*, not its slack. Discounting it gives margin away to customers who are already coming.
- **A seasonal "summer iced" push or a weather-data subscription.** Cold-drink share stays at 60.2–60.9% as Beirut warms from 14.0 °C to 25.5 °C. No weather, dust or holiday effect survives controls.
- **Winter or holiday promotions.** January is the weakest month, but 99.2% of the growth to June came from more baskets per day, not bigger baskets, and holidays show no measurable effect.
- **Bundle discounts based on how often items are bought together.** Pairs bought together don't show that a discount would add sales, and bigger baskets carry discounts by chance.

Evidence: the reports feature_loop, external_factors, executive_summary, outside_in_chains and monte_carlo, plus the discount study in analysis/outputs; all tables are in reports/tables. Effects are associations measured in historical data; financial figures are scenarios before implementation costs, not forecasts.